



Operating Systems / Sistemas de Operação

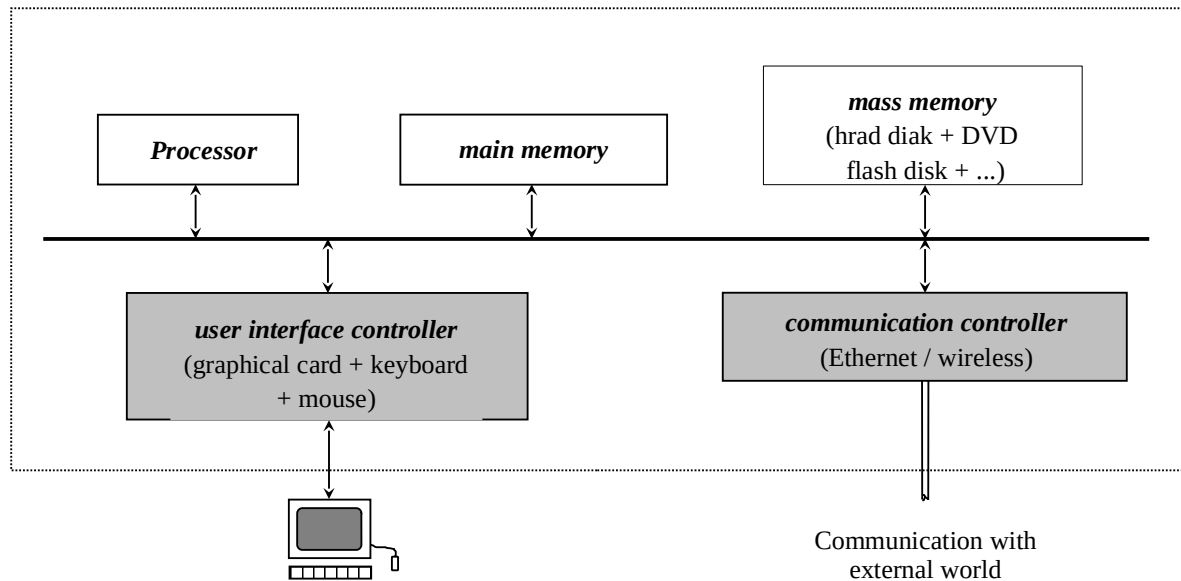
Input / Output

Artur Pereira / António Rui Borges

Outline

- *Role of the operating system*
- *Devices*
 - *Internal structure*
 - *A taxonomy*
- *Controllers*
 - *Interface with the processor*
 - *Address modes*
- *Goal of the I/O programming*
- *Types of access*
 - *Polled I/O*
 - *Interrupt-driven I/O*
 - *DMA-based I/O*

Overview



5 - 3

DETI, december 2022


Role of the operating system

- ♦ Two distinct perspectives are usually considered **for the role played by the operating system in managing the input/output devices** of a computational system
 - ♦ **user perspective** – providing the application developer with a device interface (API) that is conceptually simple, reasonably uniform, and as much as possible independent of the specific device
 - ♦ **system perspective** – isolating the different devices from direct access by user processes by introducing a functional layer that directly controls the devices
 - ♦ send commands, transfer data, handle interrupts, and handle error conditions

5 - 4

DETI, december 2022

Interface with the processor




- There are 2 different components to consider
 - *device itself* – physical system (electromechanical, optical-mechanical, ...), which stores information and converts it from, or to, an externally accessible form
 - *device controller* – electronic circuit, more or less complex, part of the computational system, which works as an interface with the device
- From the point of view of the operating system, the controller is the only relevant component
- Nowadays, controllers are very versatile, minimizing the role of the operating system in its management (programming)

Types of devices

- In terms of transferring information, input/output devices fall into two broad categories
 - *character-type devices* – the transfer of information is based on a stream of bytes, whose length can be variable
 - *block-type devices* – the transfer of information is based on a constant and pre-defined number of bytes, the *block*, typically with a value equal to a power of 2 between 512Bytes and 16KBytes
- The way the transfer is done depends on the bus used
 - bytes (8 bits), 2-bytes (16 bits), 4-bytes (32 bits) or 8-bytes (64 bits).
- The rate of transfer depends on the type of device
 - can vary from tens of bytes (keyboard, for example) to thousands of megabytes (SATA or USB3 disk)

Device controller - 1

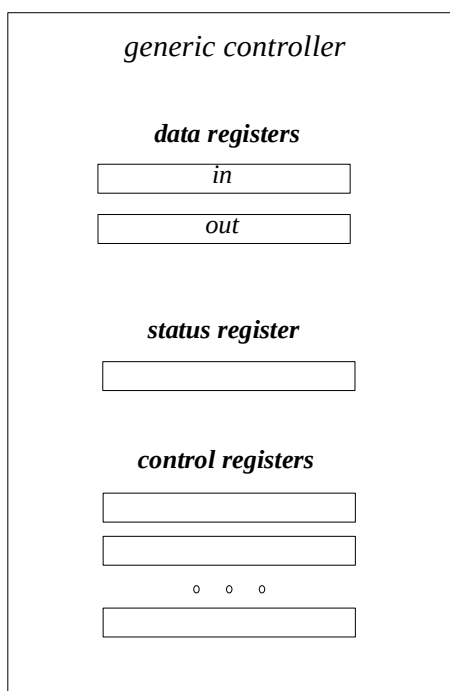


- A *generic controller*, from a programming point of view, can be seen as set of registers
- *control registers* – playing different functions
 - to configure the device
 - to define the type of interaction with the processor (*polled I/O*, *interrupt-driven I/O* or *DMA-based I/O*)
 - in complex controllers, to execute a command

5 - 7

DETI, december 2022

Device controller - 2

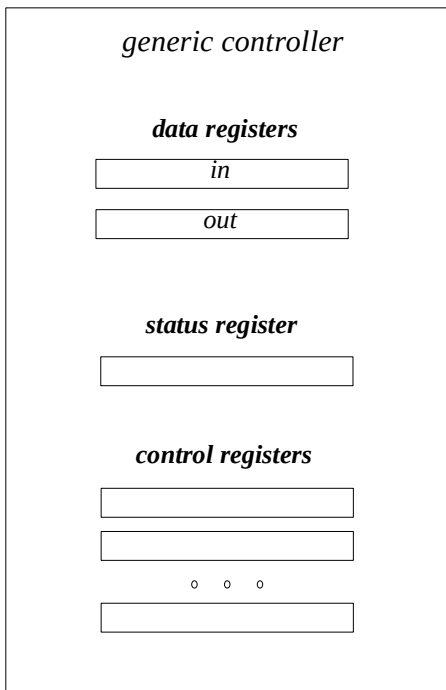


- *Status register* – representing the internal state of the device
 - to indicate the success of the last operation
 - to indicate the failure and errors of the last operation
 - to indicate it is ready to receive a new command
 - ...
- *Data registers* – used for the communication itself
 - values written in the *out* register are sent to the device
 - values read from the *in* register came from the device

5 - 8

DETI, december 2022

Device controller - 3



- In *character-type* devices, the *write* and *read* commands are implicit
 - a value written in the *out* register is sent to the device
 - a value received from the device is put in the *in* register
- In *block-type* devices, the transfer starts based on an explicit command
 - the data register is in general unique, *in-out*, and the direction of the transfer depends on the command given

5 - 9

DETI, december 2022


I/O address modes - 1

- There 3 different possible ways to access the internal registers of a controller
 - *I/O-mapped* – controllers are mapped in a specific I/O address space
 - registers are accessed through specific instructions (*in* e *out*)
 - *memory-mapped I/O* – controllers are mapped in the memory address space
 - registers are accessed through the memory access instructions (*load* e *store*)
 - *hybrid* – controllers are mapped in a specific I/O address space, but data buffers are mapped in the memory address space to facilitate communication

5 - 10


DETI, december 2022

I/O address modes - 1




- There are 3 different possible ways to access the internal registers of a controller
 - *I/O-mapped* – controllers are mapped in a specific I/O address space
 - registers are accessed through specific instructions (*in* e *out*)
 - *memory-mapped I/O* – controllers are mapped in the memory address space
 - registers are accessed through the memory access instructions (*load* e *store*)
 - *hybrid* – controllers are mapped in a specific I/O address space, but data buffers are mapped in the memory address space to facilitate communication

I/O address modes - 2



- The Intel *Pentium* has an I/O address space of 64 KB
 - Some computer systems, based on *Pentium*, use this space to address controllers
 - But, a region in memory, between addresses 640 KB e 1 MB, is also reserved to implement data buffers for the devices

I/O address modes - 3



- On a simple configuration, all resources of a computer system are connected using a single bus
- However, nowadays, personal computer systems use a broadband bus dedicated to the transfer of data between the processor and the main memory, to take advantage of the clock frequency
 - In such cases, the memory is *dual-port* to allow concurrency with the transfer of data of DMA-based devices
- Also, not all controllers are connected to the same bus
 - For example, the ISA and PCI buses can be supported by the same computer

Objectives of I/O Programming - 1

- The environment provided by the operating system for the communication with the *I/O devices* should:
 - *be independent of the device specifics*
 - devices must be seen in a generic way
 - *I/O redirecting*, for example, should be possible in a natural way
 - *support a uniform naming mechanism*
 - device names must consist of strings of characters without any particular meaning
 - *decouple devices from user processes*
 - vast majority of I/O devices work in an asynchronous manner – data transfers to and from main memory are triggered by interrupts
 - from the user's perspective, however, it is simpler to design communication in a synchronous way – the process blocks until conditions are met for communication to take place


Objectives of I/O Programming - 2

- The environment provided by the operating system for the communication with the *I/O devices* should (cont.)
 - *manage access to preemptable and non-preemptable devices in a uniform way*
 - communication with preemptible devices can be shared by multiple users simultaneously
 - communication with non-preemptible devices must take place in a mutual exclusion, or dedicated regime
 - the operating system therefore has to identify the different situations and ensure proper coordination
 - *perform error management in an integrated manner*
 - the detection of errors must be carried out as close to the device as possible in order to allow its [possible] recovery in a transparent way
 - the general policy should be to only report the error to the upper layer if the lower layer cannot handle it

5 - 15

DETI, december 2022

Polled I/O - 1




- In a *polled-I/O* approach there is no decoupling
 - It is the user process that is directly responsible for the communication
 - *Device communication routines* are *system calls* that directly implement hardware access.
- The *simplest solution*, but *little efficient*
 - The user processor enters a *busy waiting*, waiting for the completion of the operation

5 - 16

DETI, december 2022

Polled I/O - 2



```
/* access routine; assumes a character-type device */
```

```
void control(unsigned short add, unsigned char prog []);
```

```
#define RXRDY    ...    /* there are data to be read */
#define TXRDY    ...    /* transmitter register is empty */
#define ERROR    ...    /* error status */
```

```
unsigned char status(unsigned short add);
```

```
unsigned char in(unsigned short add);
```

```
void out(unsigned short add, unsigned char val);
```

Polled I/O - 3

```
/* possible routines for interaction with the device
   system calls - running at kernel level
   It is assumed that the communication channel was established already */
```

```
/**
 * \brief read N bytes
 * \param dd    --- device descriptor, which represents the communication channel
 * \param N     --- number of bytes to be read
 * \param buff  --- pointer to storage area
 * \return 0 on success; -1 on error (errno is set accordingly)
 */
```

```
int readNBytes (int dd, int N, unsigned char buff[]);
```

```
/**
 * \brief write N bytes
 * \param dd    --- device descriptor, which represents the communication channel
 * \param N     --- number of bytes to be written
 * \param buff  --- pointer to storage area
 * \return 0 on success; -1 on error (errno is set accordingly)
 */
```

```
int writeNBytes (int dd, int N, unsigned char buff[]);
```

Polled I/O - 4

```
int readNBytes(int dd, int N, unsigned char buff[])
{
    int add = getAdd(dd);

    for (int n = 0; n < N; n++)
    {
        int stat;
        do
        {
            stat = status(add);
            if ((stat & ERROR) != 0)
                /* error handling */
        } while ((stat & RXRDY) == 0);

        buff[n] = in(add);
    }
    return 0;
}
```

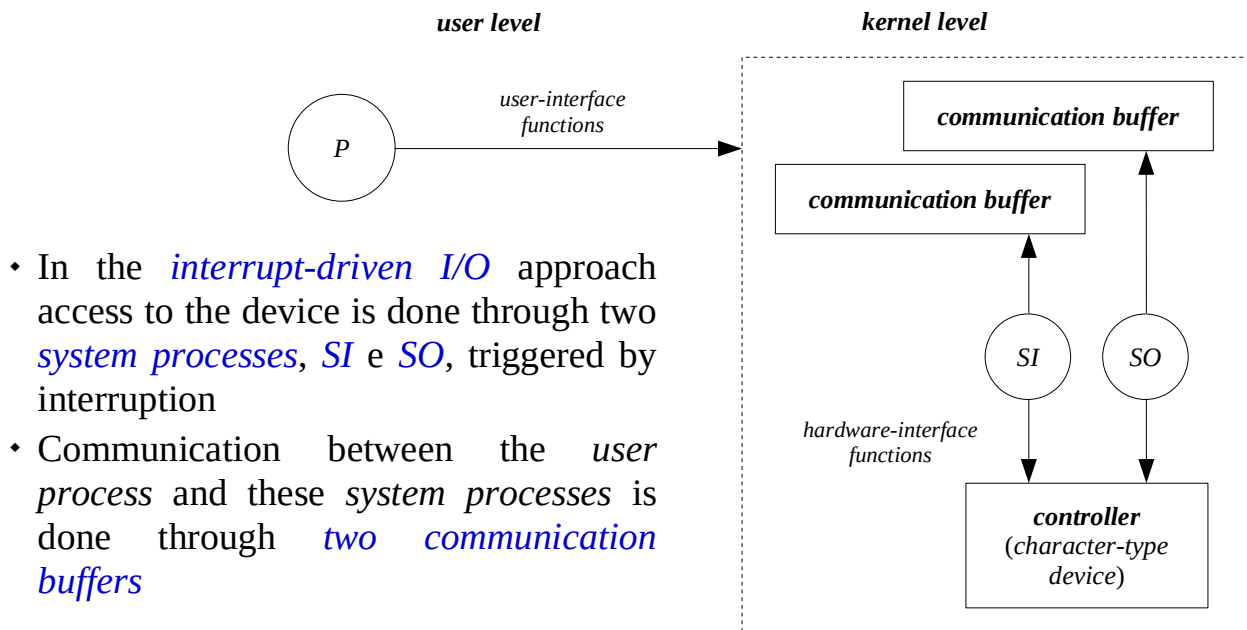
Polled I/O - 5

```
int writeNBytes(int dd, int N, unsigned char buff [])
{
    int add = getAdd(dd);

    for (int n = 0; n < N; n++)
    {
        int stat;
        do
        {
            stat = status(add);
            if ((stat & ERROR) != 0)
                /* error handling */
        } while ((stat & TXRDY) == 0);

        out(add, buff[n]);
    }
    Return 0;
}
```

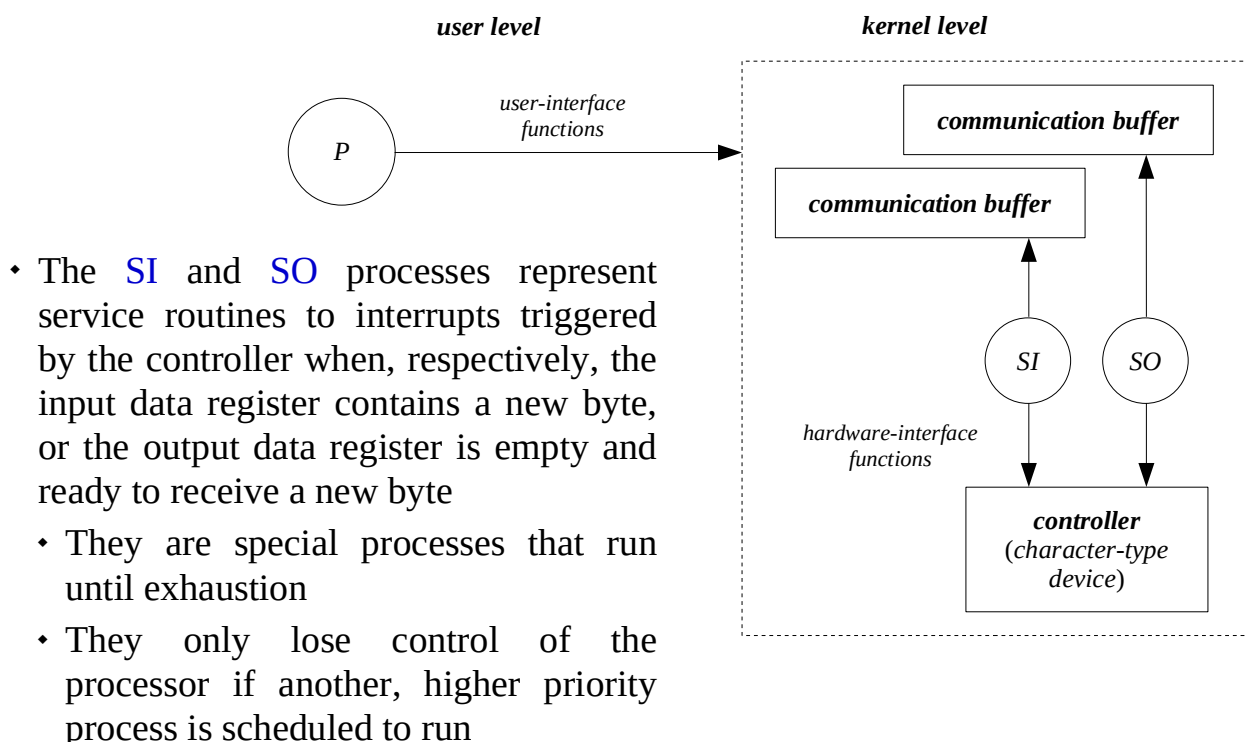
Interrupt driven I/O - 1



5 - 21

DETI, december 2022

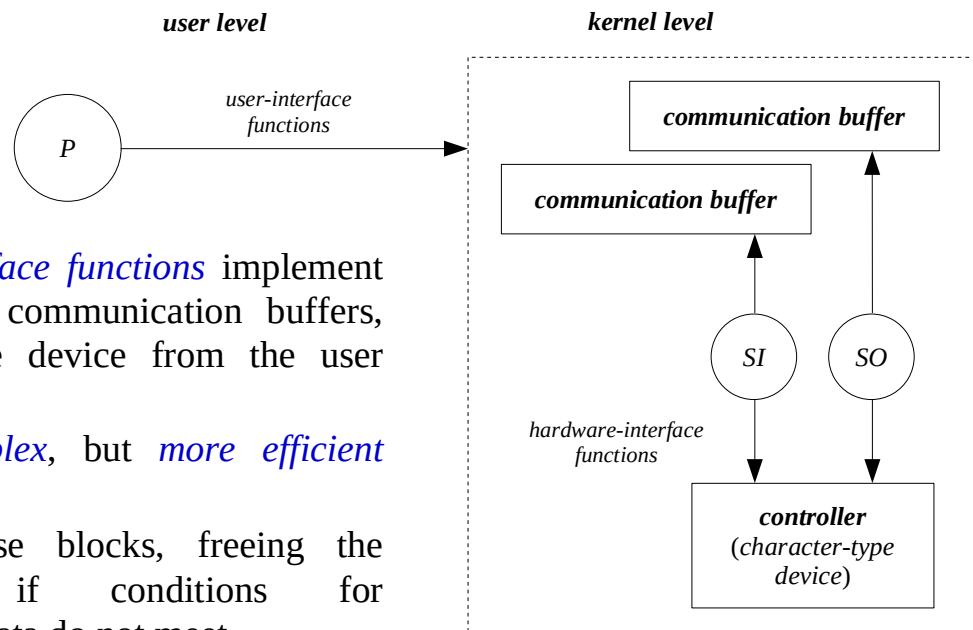
Interrupt driven I/O - 2



5 - 22

DETI, december 2022

Interrupt driven I/O - 3



- The *user-interface functions* implement access to the communication buffers, *decoupling* the device from the user process
- A *more complex*, but *more efficient* solution
 - User processes block, freeing the processor, if conditions for transferring data do not meet

Interrupt driven I/O - 2

- The SI and SO processes represent service routines to interrupts triggered by the controller when, respectively, the input data register contains a new byte, or the output data register is empty and ready to receive a new byte
 - They are special processes that run until exhaustion
 - They only lose control of the processor if another, higher priority process is scheduled to run
- The *user-interface functions* implement access to the communication buffers, *decoupling* the device from the user process
- A *more complex*, but *more efficient* solution
 - *User processes* block, freeing the processor, if conditions for transferring data do not meet

Interrupt-driven I/O - 4

```
/* access routine; assumes a character-type device */
```

```
typedef struct COM_BUFF
{
    FIFO fifo;           /* storage area */
    SEMAPHORE wait;      /* user process' blocking semaphore */
    bool noMoreInt;      /* end of interruptions (only for output) */
    int errnum;          /* error status */
} COM_BUFF;
```

Observações

- in an *entry buffer*, the semaphore is initialized at 0
 - meaning that the FIFO is empty, so no data can be read
- in an *output buffer*, the semaphore is initialized at N, the size of the fifo
 - meaning that the FIFO is empty, so N bytes can put there
- *flag* noMoreInt signals the need to prime the controller output data register so that interrupts are generated again
 - initialized to true

Interrupt-driven I/O - 5

```
/* possible routines for interaction with the device
   system calls – running at kernel level
   It is assumed that the communication channel was established already */
```

```
/**
 * \brief read N bytes
 * \param dd --- device descriptor, which represents the communication channel
 * \param N --- number of bytes to be read
 * \param buff --- pointer to storage area
 * \return 0 on success; -1 on error (errno is set accordingly)
 */
```

```
int readNBytes (int dd, int N, unsigned char buff[]);
```

```
/**
 * \brief write N bytes
 * \param dd --- device descriptor, which represents the communication channel
 * \param N --- number of bytes to be written
 * \param buff --- pointer to storage area
 * \return 0 on success; -1 on error (errno is set accordingly)
 */
```

```
int writeNBytes (int dd, int N, unsigned char buff[]);
```

Interrupt-driven I/O - 6

```

int readNBytes(int dd, int N, unsigned char buff[])
{
    COMM_BUF * inbuf = getBuff(dd);

    for (int n = 0; n < N; n++)
    {
        if (inbuf->errNumb != 0)
            /* error handling */

        sem_down(inbuf->wait);

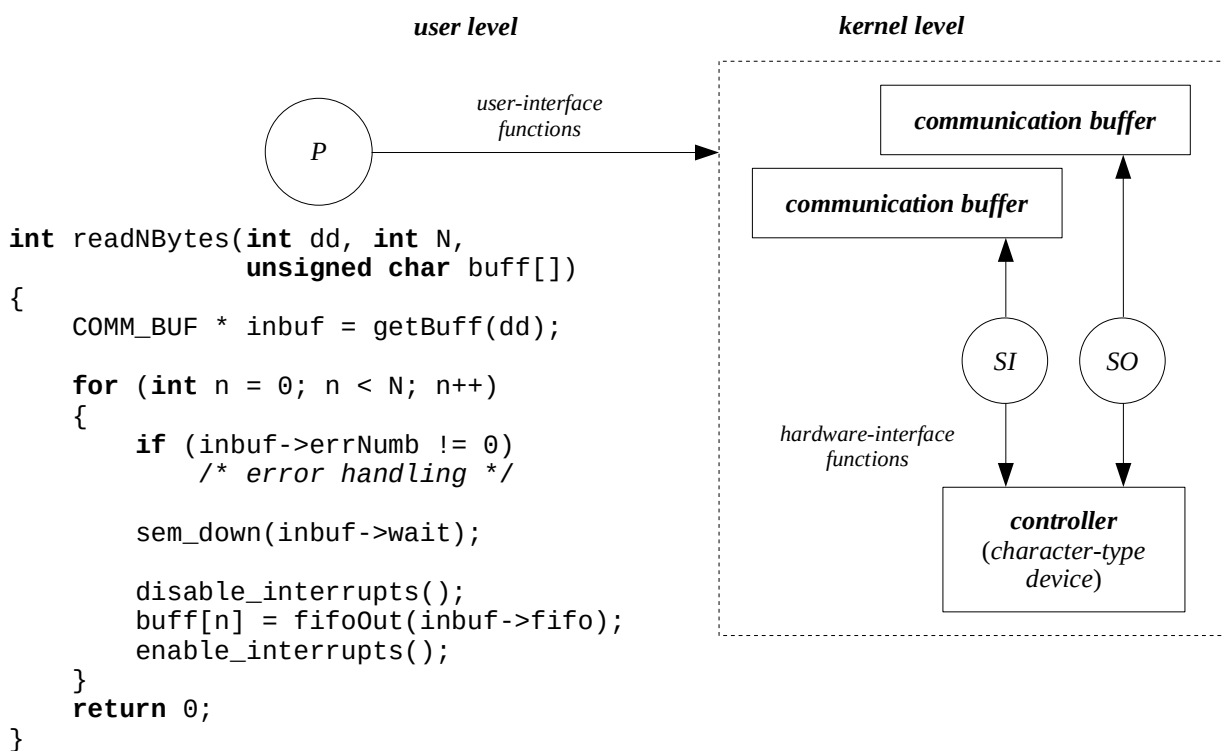
        disable_interrupts();
        buff[n] = fifoOut(inbuf->fifo);
        enable_interrupts();
    }
    return 0;
}

```

5 - 27

DETI, december 2022

Interrupt driven I/O - 6



5 - 28

DETI, december 2022

Interrupt-driven I/O - 7

```
int writeNBytes(int dd, int N, unsigned char buff [])
{
    COMM_BUF * outbuf = getBuff(dd);

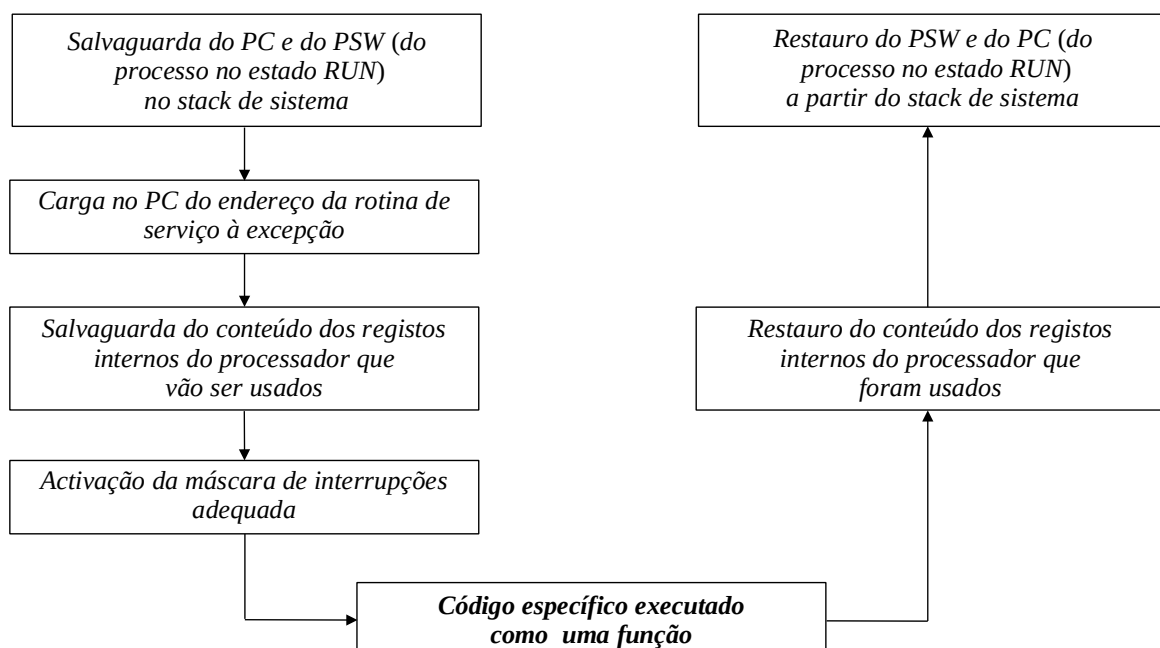
    for (int n = 0; n < N; n++)
    {
        if (outbuf->errNumb != 0)
            /* error handling */

        sem_down(outbuf->wait);

        disable_interrupts();
        fifoIn(outbuf->fifo, buff[n]);
        if (outbuf->noMoreInt)
        {
            writeDReg(getAdd(dd));
            Outbuf->noMoreInt = false;
        }
        enable_interrupts();
    }
    return 0;
}
```

Interrupt driven I/O - 7

Processo ativado por interrupção – diagrama geral de processamento



Interrupt-driven I/O - 8

```
/* SI process - reading of one byte from the controller input register
   add --- controller address */
```

```
void readDReg(unsigned short add)
{
    COM_BUFF * inbuf = getBuff(getDesc(add));

    int stat = status(add);          /* read the status register */
    if ((stat & ERROR) == ERROR)
        inBuf->errnumb = stat & ERROR; /* set error status */
    if ((stat & RXRDY) == RXRDY)      /* there are data to be read */
    {
        char val = in(add);
        if ((stat & ERROR) != ERROR)
        {
            if (not fifo_full(inbuf->fifo))
            {
                fifo_in(inbuf->fifo, val); /* store byte in buffer */
                sem_up(inbuf->wait);      /* signal there is data in buffer */
            }
            else
                inbuf->errNumb = OVERRUN; /* overrun error */
        }
    }
}
```

5 - 31

DETI, december 2022

Interrupt-driven I/O - 9

```
/* SO process - write one byte to the controller output register
   add --- controller address */
```

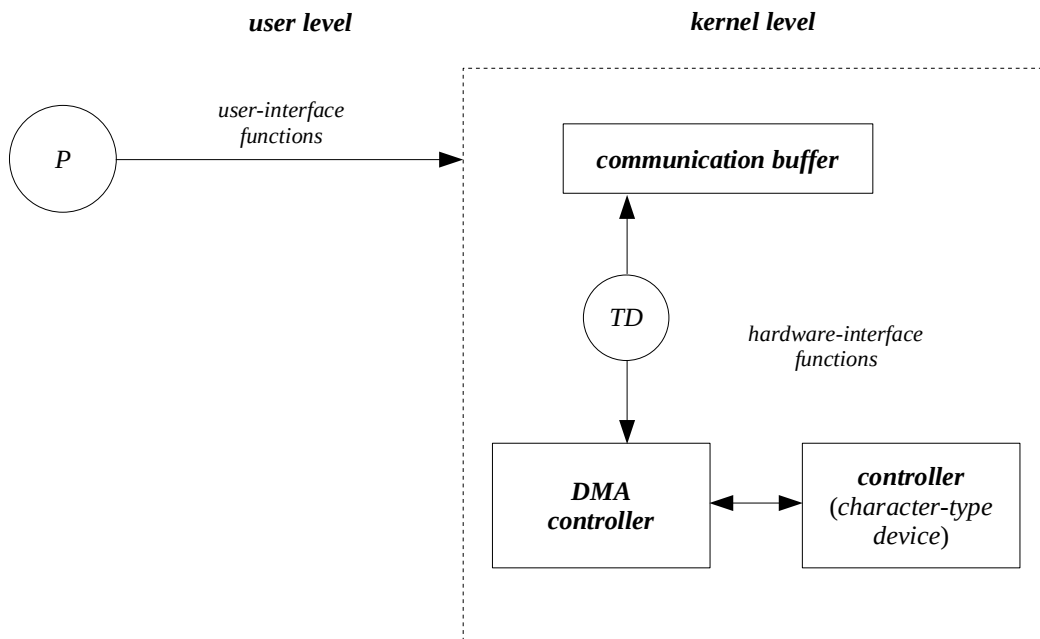
```
void writeDReg(unsigned short add)
{
    COM_BUFF * outbuf = getBuff(getDesc(add));

    int stat = status(add);          /* read status register */
    if ((stat & ERROR) == ERROR)
        outbuf->errnumb = stat & ERROR; /* set error status */
    if ((stat & TXRDY) == TXRDY)      /* can write ? */
    {
        if (not fifo_empty(outbuf->fifo))
        {
            char val = fifo_out(outbuf->fifo); /* retrieve byte from buffer */
            sem_up(outbuf->wait); /* signal there is room in buffer */
            out(add, val); /* put byte in controller output register */
        }
        else
            outbuf->noMoreInt = true; /* set end of interruptions */
    }
}
```

5 - 32

DETI, december 2022

DMA-based I/O - 1



DMA-based I/O - 2

- In a *DMA-based I/O* strategy, the DMA controller is directly connected to the device controller, or is part of it
- The underlying principle is that when the device controller wants to transfer data, it activates a *request transfer* input in the DMA controller.
- As a result, it takes control of the bus and performs one of two operations.:
 - reading a byte or word from the device controller data register and subsequently writing that value to memory (data input)
 - reading a byte or word from memory and subsequently writing that value to the device controller data register (data output)